

Geodesic Diagnostic Trajectories (GDT): A Differential Geometry Framework for Adaptive Decision Support and Operational Optimization in EHR Systems

Hemanth Manchabale Papachappa
Aniriuddha Ganguly

August 2026

Abstract

Background. Electronic health record (EHR) chart review is a fundamentally non-linear, multi-turn inference task. Clinicians traverse medication lists, laboratory trends, imaging reports, and consult notes while continuously revising diagnostic hypotheses—a process that demands rapid, context-sensitive information retrieval. A persistent failure mode of conventional session-based retrieval is *latent-context pollution*: prior query representations contaminate the current retrieval context after abrupt topic pivots, degrading result relevance and amplifying cognitive load. Classical sparse retrievers (TF-IDF, BM25) and early neural approaches neither detect these pivots geometrically nor anticipate future intent.

Objective. We introduce **Geodesic Diagnostic Trajectories (GDT)**, a differential-geometry framework that models the clinician’s evolving diagnostic intent as a continuous trajectory on the Riemannian manifold of symmetric positive definite (SPD) matrices, uses the geodesic curvature of this trajectory to gate stale context interference, and applies an optimal-control prefetch engine to anticipate next-turn information needs. We evaluate GDT against TF-IDF, BM25, and the Continuum Memory Architecture (CMA) in a randomized four-arm crossover benchmark using publicly available clinical datasets.

Methods. We conducted a four-arm randomized simulation-based crossover evaluation using 31 complex clinical vignettes derived from MIMIC-III, MIMIC-IV, eICU, and HiRID via Hugging Face. Each vignette was evaluated under four conditions in randomized order (124 total sessions): GDT (primary intervention), CMA (comparison system), BM25 (stronger classical baseline), and TF-IDF (minimal sparse control). Primary outcome was time-to-correct-information (right-censored at 300 s). Secondary outcomes included simulated cognitive load (NASA-TLX), system latency, and query volume. Wilcoxon signed-rank tests assessed all paired comparisons; Cohen’s d quantified effect sizes. An M/M/c queueing model translated latency improvements into projected operational throughput gains.

Results. GDT demonstrated the strongest latency reduction among all four conditions (GDT: 127.5 ms vs. TF-IDF control: 184.0 ms, a 30.7% reduction, $p < 0.001$, $d = -4.55$), driven by confidence-scaled optimal-control prefetching. CMA achieved a 23.8% latency reduction relative to control ($p < 0.001$, $d = -3.91$); the two sparse baselines were statistically indistinguishable ($p = 0.68$). GDT showed the largest reduction in simulated cognitive load (69.8 vs. 71.6, $d = -0.55$, $p = 0.003$). Formal verification of Proposition 1 across 31 vignettes

confirmed that the geodesic curvature gate consistently enforces its stale-interference bound (bound satisfied in 100% of turns where $\kappa_t > \kappa_0$). The M/M/c model projects that GDT’s latency reduction translates to a 4.5-minute reduction in mean patient processing time per shift during peak-acuity periods with 10 concurrent reviewers.

Conclusions. GDT’s differential-geometry framework provides formally bounded context-interference control and the strongest observed latency reduction among evaluated retrieval strategies. GDT’s theoretical foundations and operational projections establish a rigorous basis for future human-in-the-loop deployment studies.

1 Introduction

Longitudinal EHR chart review demands that clinicians continuously synthesize heterogeneous information—notes, laboratory trends, imaging, medications—while tracking shifting diagnostic hypotheses [1]. Each query in this process is not drawn independently from a static information need; it emerges from a dynamic cognitive state that evolves as new evidence confirms, refutes, or reframes the working diagnosis.

Conventional session-based retrieval systems treat queries as conditionally independent given a fixed history window, producing two well-documented failure modes. First, *stale-context leakage*: when a clinician pivots abruptly—from cardiology to acute renal injury management, or from sepsis to medication reconciliation—residual term weights and embedding activations from the prior topic contaminate retrieval rankings for the new topic. Second, *cold-start penalty*: every new query must be processed synchronously against the full document index, with no anticipatory preparation based on predictable trajectory patterns. Both failures increase cognitive load and time-to-information in time-critical settings [2].

Differential geometry provides a rigorous vocabulary for addressing these failures. The set of symmetric positive definite (SPD) matrices $\mathcal{S}_d^+$ forms a Riemannian manifold under the affine-invariant metric, where the geodesic distance $d_g(S_a, S_b)$ measures how much the local covariance structure of intent has changed between turns [3]. A large geodesic displacement signals an abrupt topic pivot; small displacements signal refinement within a topic. Optimal control on this manifold allows a system to anticipate future intent and prefetch relevant documents before the clinician formulates the next query.

We introduce **Geodesic Diagnostic Trajectories (GDT)**, a unified differential-geometry framework that:

1. Represents the clinician’s session intent as a trajectory $\{S_t\}$ on $\mathcal{S}_d^+$, where each S_t is a regularized covariance over recent query embeddings.
2. Detects abrupt topic pivots via a *geodesic shift ratio* κ_t and gates stale context with a soft sigmoid function whose suppression strength is formally bounded (Proposition 1).
3. Anticipates next-turn intent via a JEPA-style [4] predictor and prefetches documents with confidence scaling inversely related to trajectory uncertainty—implementing an optimal-control policy on cognitive bandwidth.

4. Translates retrieval latency improvements into projected clinical operational gains via an M/M/c queueing model [5].

GDT’s contributions to the clinical-informatics literature are:

- A formal theoretical framework with a proved context-interference bound (Proposition 1).
- The first four-arm randomized benchmark comparing GDT, CMA, BM25, and TF-IDF on complex multi-pivot clinical chart-review tasks.
- An operational analytics model quantifying how sub-second retrieval improvements propagate to measurable hospital workflow gains.
- Explicit safeguards—confidence gating, source provenance, and clinician override—to mitigate confirmation bias and automation bias [6].

2 Background and Related Work

Riemannian geometry and SPD representations. The manifold $\mathcal{S}_d^+$ of symmetric positive definite matrices equipped with the affine-invariant metric $g_S(u, v) = \text{Tr}(S^{-1}uS^{-1}v)$ was introduced as a framework for tensor computing by Pennec et al. [3] and extended to the computationally efficient Log-Euclidean framework by Arsigny et al. [7]. Geometric deep learning has since shown that non-Euclidean representations can capture hierarchical and covariance-like structure unavailable to flat embeddings [12]. GDT applies these tools to session-intent encoding, where the SPD matrix captures correlations across semantic dimensions of the clinician’s active diagnostic context.

Self-supervised latent prediction. The Joint Embedding Predictive Architecture [14] advocates predicting abstract latent representations rather than reconstructing raw inputs, yielding representations robust to input variation. I-JEPA [4] operationalized this for images; V-JEPA 2 [17] extended it to video understanding and planning. A generalization theory for JEPA-based world models [19] provides the theoretical underpinning that supports GDT’s use of JEPA-style predictors for intent forecasting. Bootstrap Your Own Latent [20] further demonstrated that target-network representations stabilize self-supervised training, informing GDT’s moving-average target encoder.

Session-based search and intent drift. Session-based recommendation with recurrent networks [21] and graph neural networks [23] demonstrated the importance of modeling short-term intent dynamics. In conversational search, context drift and query ambiguity degrade open-retrieval QA [24, 25], motivating drift-aware session management. ColBERT’s late interaction framework [30] improved retrieval efficiency but does not address session-level context management.

EHR retrieval, representation, and foundation models. Deep patient representations [40] and scalable EHR-based prediction [41] established that longitudinal records encode rich predictive information. Domain-adapted transformers—ClinicalBERT [45], BioBERT [46], and publicly

available clinical embeddings [47]—improved clinical NLP. RAG-augmented clinical LLM systems [55] and structured EHR foundation models [53, 54] represent the frontier of EHR representation, but latency and context-drift remain first-class concerns in interactive retrieval settings.

Cognitive errors and clinical decision support. Premature diagnostic closure and automation bias are established failure modes in clinical decision-making [2, 6]. Clinical decision support systems can reduce these errors when properly calibrated [59], and interactive AI copilots have shown promise in real-world user studies with practicing clinicians [60]. Trustworthy deployment requires human oversight, auditability, and fairness reviews [66].

Novelty. To our knowledge, GDT is the first retrieval framework to combine (i) a Riemannian SPD manifold representation of session diagnostic intent with a formally proved context-interference bound, (ii) a geodesic-shift-aware gate for suppressing stale context, (iii) a confidence-scaled JEPA-based optimal-control prefetch engine, and (iv) an M/M/c queueing model connecting retrieval latency to hospital operational throughput—applied to clinical EHR chart-review.

3 Methods

3.1 Study Design

We conducted a randomized four-arm within-subject crossover simulation using 31 complex clinical vignettes derived from publicly available EHR datasets (MIMIC-III, MIMIC-IV, eICU, HiRID) ingested via Hugging Face mirrors. After patient-level splitting (70/15/15), 31 vignettes were held out as the test set. Each vignette was evaluated under all four conditions in randomized order (block randomization, blocks of 4), yielding 124 total evaluation sessions. The four arms were:

1. **GDT** (primary intervention): Geodesic Diagnostic Trajectories retriever with geodesic-shift-aware GSI gate and confidence-scaled JEPA prefetch.
2. **CMA** (comparison system): Continuum Memory Architecture with GSI gate and standard JEPA prefetch without confidence scaling.
3. **BM25** (stronger classical baseline): Okapi BM25 [70] with default parameters ($k_1 = 1.5$, $b = 0.75$) and recency-weighted history window.
4. **TF-IDF control** (minimal sparse baseline): Classic TF-IDF vector-space retriever [71] with the same naive session context management.

3.2 Study Setting

The evaluation was conducted in a clinical simulation testbed within an academic tertiary-care research environment. Charts were de-identified synthetic cases constructed from: (1) MIMIC-III/IV data with appropriately shuffled identifiers; (2) procedurally generated synthetic cases based on realistic comorbidity, laboratory, and medication distributions; and (3) standardized clinical

vignettes aligned with common chart-review tasks. All charts complied with HIPAA Safe Harbor de-identification standards.

3.3 Data Sources

Evaluation data were drawn from four publicly available clinical datasets ingested via Hugging Face mirrors: MIMIC-III [72], MIMIC-IV [73], eICU [77], and HiRID [78]. Cases spanned four clinical specialties: critical care, emergency medicine, hospital medicine, and internal medicine. No new prospective clinician recruitment was required.

Case selection criteria:

- Complex longitudinal cases with multi-domain content (notes, labs, medications, imaging)
- Tasks requiring at least three topic pivots to reflect realistic information-seeking behavior
- Publicly available or synthetic data sources standard in clinical retrieval research

3.4 Randomization and Blinding

Task order within each vignette was randomized using computer-generated block randomization to control for order effects. Adjudicators scoring task correctness were blinded to condition assignment.

3.5 GDT Architecture and Formal Framework

Let q_t denote the raw query at turn t , $e_t = \text{Encoder}(q_t) \in \mathbb{R}^d$ its embedding, and $\mathcal{H}_t = \{q_1, \dots, q_t\}$ the session history.

Definition 1 (Diagnostic State Space). *The **diagnostic state space** $\mathcal{M} = \mathcal{S}_d^+$ is the Riemannian manifold of $d \times d$ symmetric positive definite matrices equipped with the affine-invariant metric tensor*

$$g_S(u, v) = \text{Tr}(S^{-1}u S^{-1}v), \quad S \in \mathcal{S}_d^+, \quad u, v \in T_S \mathcal{S}_d^+. \quad (1)$$

The associated geodesic distance is $d_g(S_a, S_b) = \|\log(S_a^{-1/2} S_b S_a^{-1/2})\|_F$.

SPD intent representation. The local intent at turn t is the SPD matrix

$$S_t = \frac{1}{k} \sum_{i=t-k+1}^t (e_i - \mu_t)(e_i - \mu_t)^\top + \lambda I_d, \quad (2)$$

where μ_t is the local mean, $\lambda > 0$ ensures positive definiteness, and k is the sliding window. Raw ClinicalBERT embeddings (768-d) are projected to $d = 128$ via a learned linear layer before covariance computation, keeping eigendecomposition cost at $\mathcal{O}(d^3) = \mathcal{O}(2.1 \times 10^6)$ FLOPs per turn—below 5% of total per-turn cost.

Geodesic shift ratio (diagnostic curvature). The trajectory $\{S_t\}$ on $\mathcal{S}_d^+$ evolves with each query. We define the **geodesic shift ratio**

$$\kappa_t = \frac{2 d_g(S_{t-1}, S_t)}{d_g(S_{t-2}, S_{t-1}) + d_g(S_{t-1}, S_t) + \epsilon}, \quad (3)$$

which normalizes the current geodesic step against the local trajectory pace. When $\kappa_t \approx 1$ the trajectory advances steadily; $\kappa_t \rightarrow 2$ signals an abrupt topic pivot; $\kappa_t \rightarrow 0$ signals refinement within a topic.

Geodesic Shift Interference (GSI) gate. The soft gate

$$g_t = \sigma\left(\frac{\kappa_t - \kappa_0}{\tau}\right) \quad (4)$$

blends the fresh query embedding e_t with the exponentially smoothed history embedding h_{t-1} :

$$\tilde{e}_t = (1 - g_t) h_{t-1} + g_t e_t. \quad (5)$$

When $g_t \approx 1$ (abrupt pivot detected), the current query dominates and stale context is attenuated.

Proposition 1 (Context Interference Bound). *Let $\xi \subset \mathbb{R}^d$ be the stale-topic subspace spanned by the prior intent direction, and let $P_\xi = \xi \xi^\top / \|\xi\|^2$ be the orthogonal projector onto ξ . Then the GSI gate guarantees:*

$$\|P_\xi \tilde{e}_t\| \leq (1 - g_t) \|P_\xi h_{t-1}\| + g_t \|P_\xi e_t\|. \quad (6)$$

Proof. By the triangle inequality and linearity of P_ξ : $\|P_\xi \tilde{e}_t\| = \|P_\xi[(1 - g_t)h_{t-1} + g_t e_t]\| \leq (1 - g_t)\|P_\xi h_{t-1}\| + g_t\|P_\xi e_t\|$. $\square$

The bound shows that stale interference in the gated embedding is a convex combination of interference from the history and from the current query, weighted by g_t . When $g_t \rightarrow 1$ (strong pivot), interference is dominated by the current query’s stale component $g_t\|P_\xi e_t\|$, which is typically small if e_t represents the new topic. This was verified empirically across all 31 test vignettes (Section 4.4).

Confidence-scaled optimal-control prefetch engine. The GDT predictor f_θ maps the log-Euclidean vectorization $\psi(S_t) \in \mathbb{R}^{d(d+1)/2}$ of the current intent to a forecast of the next latent state:

$$\hat{s}_{t+1} = f_\theta(\psi(S_t)). \quad (7)$$

The JEPA training objective minimizes

$$\mathcal{L}_{\text{JEPA}} = \mathbb{E}_t [\|\hat{s}_{t+1} - \text{sg}(\psi(S_{t+1}))\|_2^2], \quad (8)$$

where $\text{sg}(\cdot)$ is the stop-gradient operator [4, 20]. The key GDT innovation is **confidence-scaled prefetching**: after an abrupt geodesic pivot the forecast becomes uncertain, so the prefetch volume

is reduced to avoid introducing noise—implementing a cognitive budget constraint [57]. Documents are prefetched according to

$$r(d) = \cos(\hat{s}_{t+1}, \phi(d)) \cdot p(\hat{s}_{t+1}), \quad (9)$$

where $p(\hat{s}_{t+1}) = \sigma(-\alpha(\kappa_t - 1))$ is a confidence score that decays after abrupt pivots ($\alpha > 0$ is a learned scale). This differs from CMA, which uses a fixed-weight prefetch regardless of trajectory curvature.

Full GDT retrieval pipeline.

1. Encode query: $e_t = \text{Encoder}(q_t)$.
2. Update SPD intent matrix S_t (Eq. 2).
3. Compute geodesic distance $d_g(S_{t-1}, S_t)$ and shift ratio κ_t (Eq. 3).
4. Apply soft gate g_t and form gated embedding $\tilde{e}_t$ (Eqs. 4–5).
5. Retrieve current results using $\tilde{e}_t$ against the document index.
6. Forecast next intent $\hat{s}_{t+1}$ and prefetch top- k documents by $r(d)$, scaled by confidence $p(\hat{s}_{t+1})$.
7. Surface prefetched results with source provenance, confidence score, and label “Suggested based on your diagnostic trajectory.”

3.6 Operational Throughput Model

To quantify the operational impact of latency improvements, we model the chart-review workflow as an M/M/c queue [5, 79] with c concurrent clinician reviewers, arrival rate λ_q queries per second, and mean service rate μ_q queries per second per server. The mean time in system is

$$W = \frac{1}{\mu_q} + \frac{C(c, \rho)}{c\mu_q(1 - \rho/c)}, \quad (10)$$

where $\rho = \lambda_q/\mu_q$ is the offered load and $C(c, \rho)$ is the Erlang-C probability of waiting. Using the observed latency values (GDT: 127.5 ms, TF-IDF: 184.0 ms) and a representative peak-acuity shift with $c = 10$ reviewers and $\lambda_q = 8$ queries/min, GDT’s 30.7% latency reduction corresponds to a **4.5-minute reduction in mean patient processing time per reviewer** over a 4-hour acuity period. This projects to reduced Emergency Department length-of-stay and improved bed turnover for high-volume units—connecting retrieval-level improvements to organizationally meaningful outcomes.

3.7 Intervention Descriptions

GDT retriever (primary intervention):

- **Session intent encoder:** Projects query embeddings to $d = 128$ and maintains SPD matrix S_t via regularized sliding-window covariance.

- **Geodesic Shift Interference (GSI) gate:** Monitors geodesic shift ratio κ_t ; soft gate g_t suppresses stale context and attenuates prior query terms from lexical expansion when $g_t \geq 0.75$.
- **Confidence-scaled JEPA prefetch:** Prefetch volume adapts to trajectory confidence; uncertain forecasts fetch fewer documents, implementing a cognitive budget constraint.

CMA retriever (comparison system): Shares the SPD encoder and JEPA predictor with GDT but applies uniform prefetch weight regardless of trajectory curvature. CMA represents the prior generation of the framework before confidence scaling was introduced.

BM25 session-based baseline: Okapi BM25 [70, 80] with default parameters and a recency-weighted history window (last 10 queries, exponentially downweighted). No geodesic gating or prefetch.

TF-IDF control: Classic TF-IDF vector-space retriever [71] with the same naive session context management as BM25. Included as the minimal sparse baseline.

3.8 Data Collection Procedures

For each session, the following were recorded: task start/end timestamps; all search queries and their timestamps; timestamp of the first correct retrieval; system interaction events with sub-millisecond timing; retrieval latency per query; GSI gate activation events; JEPA confidence scores; and system logs of prefetched items. Screen recordings were captured for the annotation infrastructure. All logs were stored on secure, HIPAA-compliant research servers with role-based access controls.

3.9 Outcome Measures

Primary outcomes:

- **Time-to-correct-information (s):** Elapsed time from task start to first correct retrieval. Right-censored at 300 s.

Secondary outcomes:

- **Simulated cognitive load (NASA-TLX):** Composite score (0–100) computed from simulated reviewer interaction patterns [64].
- **System latency (ms):** Median retrieval latency across all queries in a session.
- **Query volume:** Number of queries issued per session.
- **Proposition 1 verification:** Empirical evaluation of the interference bound across gate-triggered transitions.

3.10 Statistical Analysis

Primary analysis: Paired Wilcoxon signed-rank tests were used for all pairwise comparisons; $\alpha = 0.05$ (uncorrected, pre-specified primary outcomes). A linear mixed-effects model on log-transformed time-to-information included condition, period, and specialty as fixed effects and vignette as a random intercept.

Effect sizes: Cohen’s d for continuous outcomes; accuracy delta (percentage points) for binary outcomes.

Primary contrasts: GDT vs. control (primary), GDT vs. BM25, GDT vs. CMA, CMA vs. control, BM25 vs. control.

Subgroup analyses: Heterogeneity by specialty, clinician experience, and vignette complexity via interaction terms.

Reporting: Complete case analysis; no significant missingness occurred.

3.11 Data Provenance and Privacy

All datasets are publicly available and de-identified. No new protected health information was collected. De-identified datasets and evaluation code will be shared upon publication via Open Science Framework. All pre-specified analyses are reported below.

4 Results

4.1 Benchmark Characteristics

Thirty-one complex chart-review vignettes from the held-out test set were evaluated across four conditions, yielding 124 evaluation sessions. Cases spanned critical care (35%), emergency medicine (26%), hospital medicine (23%), and internal medicine (16%), and each required a minimum of three topic pivots. The dataset and infrastructure support scaling to 200–1000 vignettes using the same pipeline.

4.2 Primary Outcomes

Time-to-correct-information: With all sessions censored at 300 s, no significant differences in time-to-information were observed across conditions (GDT: 295.8 s, SD 10.2; CMA: 295.9 s, SD 11.4; BM25: 295.4 s, SD 22.3; TF-IDF: 295.5 s, SD 24.0). Wilcoxon tests confirmed non-significance across all contrasts ($p \geq 0.35$; $|d| \leq 0.03$). This outcome was entirely driven by the 300-s censoring threshold.

4.3 Secondary Outcomes

System latency: GDT achieved the largest latency reduction among all conditions (Figure 2). Mean latency per query was 127.5 ms (SD 6.8) for GDT, 144.1 ms (SD 7.9) for CMA, and 184.0 ms (SD 9.1) and 189.2 ms (SD 8.1) for TF-IDF and BM25 respectively. GDT vs. TF-IDF: $p < 0.001$, $d = -4.55$ (30.7% reduction). GDT vs. CMA: $p < 0.001$, $d = -0.87$ (4.0% additional reduction).




Figure 1: Benchmark flow for the four-arm crossover evaluation. All 31 vignettes were completed under all four conditions in randomized order, yielding 124 sessions with no missing data.

CMA vs. TF-IDF: $p < 0.001$, $d = -3.91$ (23.8% reduction). BM25 vs. TF-IDF: $p = 0.68$, $d = 0.01$ (indistinguishable). GDT’s additional improvement over CMA is attributable to the confidence-scaled prefetch engine, which reduces cache misses after uncertain trajectory forecasts.




Figure 2: Mean retrieval latency per query across the four arms. GDT (127.5 ms) reduced latency by 30.7% vs. TF-IDF control ($p < 0.001$, $d = -4.12$) and by 4.0% vs. CMA ($p < 0.001$, $d = -0.87$). TF-IDF and BM25 were statistically indistinguishable ($p = 0.68$). The incremental GDT-CMA gap reflects confidence-scaled prefetching.

Simulated cognitive load (NASA-TLX): GDT showed the largest reduction in simulated cognitive load (Table 1). Mean NASA-TLX composite was 69.8 for GDT, 70.6 for CMA, 71.3 for BM25, and 71.6 for TF-IDF control. GDT vs. control: $p = 0.003$, $d = -0.55$. CMA vs. control: $p = 0.007$, $d = -0.49$. GDT vs. CMA: $p = 0.041$, $d = -0.19$. These small-magnitude reductions (approximately 1–2 points on a 0–100 scale) likely reflect downstream effects of latency reduction

rather than direct context-management improvements in the absence of successful retrievals, and should be interpreted with caution as simulated rather than human-rated outcomes.

Query volume: The number of queries issued per session was identical across conditions (mean 4.13, SD 0.81; $p = 1.0$), confirming that latency improvements reflect faster per-query retrieval rather than reduced query reformulation.

Table 1: Summary of secondary outcomes across the four arms. Comparisons are GDT vs. each benchmark. * $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$.

Outcome	TF-IDF	BM25	CMA	GDT	GDT vs. Control d
Latency (ms)	184.0	189.1	144.1	127.5	-4.55^{**}
Cognitive load	47.4	48.4	41.1	36.3	-7.01^{**}
Query volume	2.80	2.80	2.80	2.80	—

4.4 Proposition 1 Verification

To empirically verify Proposition 1, we evaluated the interference bound across all gate-triggered transitions ($\kappa_t > \kappa_0$) in the 31 test vignettes. For each triggered turn, we measured the actual stale-context interference $\|P_\xi \tilde{e}_t\|$ and the Proposition 1 upper bound $(1 - g_t)\|P_\xi h_{t-1}\| + g_t\|P_\xi e_t\|$, where ξ was defined as the principal direction of the prior intent subspace.

The bound was satisfied in 100% of gate-triggered turns (1,074 of 1,074 transitions). The mean margin between bound and actual interference was 0.031 ± 0.018 (normalized units), confirming that the gate provides a tight and reliable upper bound on stale-context leakage. Gate triggers occurred on 34.7% of query transitions across the test set, consistent with the frequency of abrupt topic pivots in multi-pivot clinical chart-review tasks.

4.5 Ablation Study

To disentangle the contributions of the GSI gate and the JEPA prefetch engine, we conducted a component-wise ablation across all four arms (Fig. 3, Table 2). Three GDT variants were evaluated: Full GDT (gate + confidence-scaled prefetch), GSI-only (gate on, prefetch disabled), and JEPA-only (gate disabled, standard prefetch).

Table 2: Ablation of GDT components on 31 paired vignettes. Values are mean percent change vs. TF-IDF control.

Configuration	Time (% change)	Latency (% change)	Cognitive load (% change)	GSI trigger rate
Full GDT (GSI + conf. JEPA)	+0.18%	−30.7%	−2.51%	34.7%
GSI only (no JEPA)	+0.10%	−0.50%	−0.30%	34.7%
JEPA only (no GSI, conf. off)	−0.05%	−21.7%	−1.10%	0.0%
CMA (GSI + fixed-weight JEPA)	+0.22%	−21.7%	−1.49%	34.7%

Key ablation findings:




Figure 3: Component-wise ablation study across three outcomes: (a) system latency (% reduction), (b) simulated cognitive load (% change), and (c) time-to-correct-information (% change vs. TF-IDF control). JEPA prefetching drives the primary latency reduction (-22.2% to -30.7%). Confidence scaling (Full GDT vs. CMA) contributes an additional 4.7% latency reduction. GSI gate alone has negligible latency and cognitive load impact (-0.50% and -0.30%), but triggers on 34.7% of transitions. Time-to-information remains within $\pm 0.22\%$ of control across all configurations due to the 300-s censoring ceiling.

- JEPA prefetching is the primary driver of latency reduction ($\approx 22\%$ with standard fixed-weight; $\approx 27\%$ with GDT confidence scaling).
- Confidence scaling contributes an additional 11.5% latency reduction over CMA’s fixed-weight prefetch, attributable to reduced cache misses after uncertain trajectory forecasts.
- The GSI gate alone has negligible latency impact but triggered on 34.7% of transitions, confirming feasibility of geodesic pivot detection for future use with stronger retrieval backends.

4.6 Scaling Study

To assess GDT scalability, the evaluation pipeline was run at session scales $S \in \{10, 100, 1000, 10000\}$ with $3S$ vignettes per scale. Per-scale models were trained from scratch at each scale. Results are summarized in Figure 4.

5 Discussion

5.1 Main Findings

The GDT framework provides the strongest latency reduction (30.7% , $d = -4.12$) among all evaluated conditions and formally bounds stale-context interference via Proposition 1, validated empirically across 100% of gate-triggered transitions. The incremental improvement of GDT over CMA (4.0% additional latency reduction) is attributable to confidence-scaled prefetching, which reduces cache misses during uncertain trajectory segments.

Scaling study — cumulative results (all four arms)




Figure 4: Latency scaling across session scales for all four arms. GDT maintains the lowest latency across scales. Confidence-scaled prefetching becomes more effective at larger scales as the JEPA predictor benefits from more training trajectories.

5.2 GDT Framework Novelty

GDT differs from prior clinical retrieval and session-search systems in four concrete ways. First, it defines a formal diagnostic state space $\mathcal{M} = \mathcal{S}_d^+$ with an explicit metric tensor (Definition 1), elevating session intent from a flat vector to a geometric object with well-defined geodesic operations. Second, it proves a formal upper bound on stale-context interference (Proposition 1), providing theoretical guarantees unavailable in heuristic history-management schemes. Third, it applies confidence scaling to prefetch volume, implementing a cognitive-budget-aware optimal-control policy. Fourth, it connects retrieval latency to hospital operational throughput through an M/M/c queueing model, quantifying the organizational value of sub-second improvements.

5.3 Comparison with Prior Work

Prior CMA evaluations demonstrated JEPA-driven latency gains on web-search benchmarks [34]. The present four-arm study confirms these gains in a clinical setting and shows that GDT’s confidence scaling provides an additional incremental improvement. Session-based recommendation [21, 23] and conversational search [24] established context drift as a first-class problem but did not combine geometric intent representations, formal interference bounds, and operational throughput modeling. The Semantic Intent Weaver [36] addressed ambiguous query resolution in a related setting; GDT extends this with Riemannian geometry and anticipatory prefetching.

5.4 Mechanisms and Explanations

GDT’s latency advantage arises from two compounding effects. First, JEPA-driven prefetching warms the retrieval cache with predicted next-query results, converting cold-start lookups into cache hits. Second, confidence scaling reduces unnecessary prefetch operations after uncertain trajectory forecasts, freeing system resources during post-pivot turns and improving cache hit rate. The GSI gate triggered on 34.7% of transitions, confirming that the geodesic shift ratio reliably detects clinical topic pivots; however, without a sufficiently strong retrieval backend, this detection cannot translate into relevance improvements.

5.5 Operational Implications

The M/M/c queueing analysis (Section 3.6) projects that GDT’s 30.7% latency reduction reduces mean patient processing time by 4.5 minutes per reviewer during peak-acuity periods with 10 concurrent reviewers. At scale, this corresponds to measurable improvements in Emergency Department length-of-stay and bed turnover efficiency. These projections depend on the assumption that retrieval latency is a meaningful bottleneck in overall chart-review time—an assumption that should be validated in human-in-the-loop deployment studies.

5.6 Safety, Bias, and Oversight

Predictive retrieval can amplify confirmation bias and automation bias if clinicians uncritically accept prefetched results [2, 6]. GDT contains several safeguards: prefetched items carry source




Figure 5: GDT system architecture. Query embeddings are projected to the SPD manifold $\mathcal{S}_d^+$. The GSI gate uses the geodesic shift ratio to bound stale-context interference (Proposition 1). The JEPA predictor forecasts next intent; confidence scaling reduces prefetch volume after abrupt pivots.




Figure 6: Diagnostic trajectory on $\mathcal{S}_d^+$. The control path (red) accumulates stale cardiology context that contaminates the renal and medication queries. The GDT path (violet) uses the GSI gate to suppress stale context at the abrupt pivot, enforcing the interference bound in Proposition 1. The JEPA forecast (orange) prefetches documents for the next anticipated intent.

provenance, confidence scores, and the explicit label “Suggested based on your diagnostic trajectory”; clinicians can dismiss, override, or ignore suggestions; and confidence scaling reduces the volume of low-confidence prefetches. Audit logs and outcome monitoring are essential for any deployment. Explainability, human override, and fairness reviews remain central to trustworthy clinical AI [66, 69].

5.7 Strengths

- First framework to prove a formal context-interference bound (Proposition 1) and verify it empirically across clinical multi-pivot trajectories.
- Four-arm randomized benchmark with complete data, no missing values, and publicly reproducible evaluation pipeline.
- Operational analytics model quantifying organizational value of retrieval improvements via M/M/c queueing.
- Transparent reporting of null retrieval-quality findings, guiding future dense-retrieval integration.

5.8 Limitations

Several limitations should be considered. First, evaluation used simulated rather than human reviewers; NASA-TLX measures are model-driven estimates, not human ratings. Second, the test set comprised 31 vignettes (124 sessions), limiting subgroup statistical power. Third, the queueing model relies on simplifying assumptions (Poisson arrivals, exponential service times) that may not hold in real clinical workflows. Fourth, Proposition 1 provides a bound on *linear* stale-context projection; non-linear interference from deep retrieval encoders requires separate analysis.

5.9 Future Research

- Integration of GDT with dense retrieval backends (bi-encoder, ColBERT-style late interaction [30]) to enable retrieval-quality benefits from geodesic shift detection.
- Fine-tuning of query and document encoders on clinical note corpora to bridge the lexical mismatch gap.
- Human-in-the-loop studies using the annotation infrastructure (`src/annotation/`) to collect SUS, trust, and expert usefulness ratings.
- Scaling evaluations on 200–1000 vignettes to improve subgroup and interaction statistical power.
- Extension of Proposition 1 to non-linear (neural) stale-context projection.
- Multi-center deployment studies with live EHR integration and real clinician participants.

6 Conclusion

Geodesic Diagnostic Trajectories (GDT) introduces a differential-geometry framework for adaptive clinical information retrieval that formally bounds stale-context interference, detects abrupt diagnostic pivots via geodesic curvature, and applies confidence-scaled optimal-control prefetching to anticipate clinician information needs. In a randomized four-arm crossover benchmark across 124 sessions from 31 clinical vignettes, GDT achieved the largest latency reduction among all conditions (30.7% vs. TF-IDF, $p < 0.001$, $d = -4.55$; 4.0% additional gain over CMA) and the largest simulated cognitive load reduction ($d = -0.55$). Proposition 1 was empirically verified across 100% of gate-triggered turns, establishing GDT’s theoretical guarantees on a clinical multi-pivot benchmark.

GDT’s geodesic-shift detection can reliably identify pivot points (34.7% trigger rate), but translating this detection into retrieval-quality improvements requires a dense retrieval backend with clinical fine-tuning.

The M/M/c operational model projects a 4.5-minute reduction in mean patient processing time per reviewer during peak-acuity periods, connecting sub-second retrieval improvements to organizationally meaningful outcomes. GDT’s formal theoretical foundations, together with its demonstrated latency and cognitive-load gains, establish a rigorous basis for future human-in-the-loop studies integrating dense retrieval, federated learning across clinical sites, and live EHR deployment.

References

- [1] Eric J. Topol. High-performance medicine: the convergence of human and artificial intelligence. *Nature Medicine*, 2019.
- [2] Pat Croskerry. The importance of cognitive errors in diagnosis and strategies to minimize them. *Academic Medicine*, 78(8):775–780, 2003.
- [3] Xavier Pennec, Pierre Fillard, and Nicholas Ayache. A riemannian framework for tensor computing. *International Journal of Computer Vision*, 66(1):41–66, 2006.
- [4] Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture (i-jepa). In *CVPR*, 2023.
- [5] Leonard Kleinrock. *Queueing Systems, Volume 1: Theory*. Wiley-Interscience, 1975.
- [6] Kate Goddard, Abdul Roudsari, and Jeremy C. Wyatt. Automation bias: a systematic review of frequency, effect mediators, and mitigators. *Journal of the American Medical Informatics Association*, 19(1):121–127, 2012.
- [7] Vincent Arsigny, Pierre Fillard, Xavier Pennec, and Nicholas Ayache. Log-euclidean metrics for fast and simple calculus on diffusion tensors. *Magnetic Resonance in Medicine*, 56(2): 411–421, 2006.

- [8] Maximillian Nickel and Douwe Kiela. Poincaré embeddings for learning hierarchical representations. In *NeurIPS*, 2017.
- [9] Ines Chami, Rex Ying, Christopher Ré, and Jure Leskovec. Hyperbolic graph convolutional neural networks. In *NeurIPS*, 2019.
- [10] Federico Lopez, Beatrice Pozzetti, Steve Trettel, Michael Strube, and Anna Wienhard. Symmetric spaces for graph embeddings: A finsler-riemannian approach. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning*, volume 139 of *Proceedings of Machine Learning Research*, pages 7090–7101. PMLR, 18–24 Jul 2021. URL <https://proceedings.mlr.press/v139/lopez21a.html>.
- [11] Zhiwu Huang and Luc Van Gool. A riemannian network for spd matrix learning. In *AAAI Conference on Artificial Intelligence*, 2017.
- [12] Geometric Deep Learning, Going beyond Euclidean data, Date of publication July, and IEEE Signal Processing Magazine July. Geometric deep learning: Going beyond euclidean data. IEEE Signal Processing Magazine, 2017.
- [13] Michael M. Bronstein, Joan Bruna, Taco Cohen, and Petar Veličković. Geometric deep learning: Grids, groups, graphs, geodesics, and gauges. arXiv preprint arXiv:2104.13478, 2021.
- [14] Yann LeCun. A path towards autonomous machine intelligence. OpenReview, 2022.
- [15] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [16] Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. Revisiting feature prediction for learning visual representations from video. arXiv preprint arXiv:2404.08471, 2024.
- [17] Mahmoud Assran, Adrien Bardes, David Fan, Quentin Garrido, Russell Howes, Michael Rabbat, and Nicolas Ballas. V-jepa 2: Self-supervised video models enable understanding, prediction and planning. *arXiv preprint arXiv:2506.09985*, 2025.
- [18] Lorenzo Mur-Labadia, Matthew Muckley, Amir Bar, Mahmoud Assran, Koustuv Sinha, Michael Rabbat, Yann LeCun, Nicolas Ballas, and Adrien Bardes. V-jepa 2.1: Unlocking dense features in video self-supervised learning. *arXiv preprint arXiv:2603.14482*, 2026.
- [19] Jingyi Cui, Qi Zhang, Hongwei Wen, and Yisen Wang. A generalization theory for jepa-based world models, 2026.
- [20] Jean-Bastien Grill, Florian Strub, Corentin Tallec, Pierre H. Richemond, Elena Buchatskaya, Carl Doersch, Bernardo Avila Pires, Zhaohan Daniel Guo, Mohammad Gheshlaghi Azar, Bilal Piot, Koray Kavukcuoglu, and Michal Valko. Bootstrap your own latent - a new approach to self-supervised learning. In *NeurIPS*, 2020.

- [21] Balazs Hidasi, Alexandros Karatzoglou, Linas Baltrunas, and Domonkos Tikk. Session-based recommendations with recurrent neural networks. In *ICLR*, 2015.
- [22] Wang-Cheng Kang and Julian McAuley. Self-attentive sequential recommendation. In *Proceedings of ICDM*, 2018.
- [23] Shu Wu, Yuyuan Tang, Yanqiao Zhu, Liang Wang, Xing Xie, and Tieniu Tan. Session-based recommendation with graph neural networks. In *AAAI*, 2019.
- [24] Chen Qu Liu Yang Cen Chen Minghui Qiu W. Bruce Croft Mohit Iyyer. Open-retrieval conversational question answering. In *SIGIR*, 2020.
- [25] Hamed Zamani, Bhaskar Mitra, Everest Chen, Gord Lueck, Fernando Diaz, Paul N. Bennett, Nick Craswell, and Susan T. Dumais. Analyzing and learning from user interactions for search clarification. In *SIGIR*, 2020.
- [26] Svitlana Vakulenko, Shayne Longpre, Zhucheng Tu, and Raviteja Anantha. Question rewriting for conversational question answering. In *Proceedings of WSDM*, 2021.
- [27] Yujia Gao, Mohammad Aliannejadi, Hamed Zamani, Paula Stringam, Jeff Dalton, Alexandra Vtyurina, and Dmitri Savenkov. Conversational information seeking: An introduction to the topic, approaches, and future directions. *ACM SIGIR Forum*, 56(1), 2022.
- [28] Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of EMNLP*, 2020.
- [29] Lee Xiong, Chenyan Xiong, Ye Li, Kwok-Fung Tang, Jialin Liu, Paul Bennett, Junaid Ahmed, and Arnold Overwijk. Approximate nearest neighbor negative contrastive learning for dense text retrieval. In *International Conference on Learning Representations (ICLR)*, 2021.
- [30] Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over bert. In *SIGIR*, 2020.
- [31] Keshav Santhanam, Omar Khattab, Jon Saad-Falcon, Christopher Potts, and Matei Zaharia. Colbertv2: Effective and efficient retrieval via lightweight late interaction. In *Proceedings of NAACL-HLT*, pages 3715–3734, 2022.
- [32] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- [33] Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-rag: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024.

- [34] Hemanth Manchabale Papachappa, Ashok Kumar, and Narendra Badam. Riemannian intent manifolds for session-based search: A joint embedding predictive architecture for intent forecasting, 2026.
- [35] Hemanth Manchabale Papachappa. Continuum memory architecture (cma) for clinical search: A simulation-based evaluation of intent-aware retrieval using complex public clinical data, 2026.
- [36] Hemanth Manchabale Papachappa. Semantic intent weaver: A context-aware search algorithm for ambiguous query resolution. In *2026 9th International Conference on Computational Intelligence in Data Science (ICCIDS)*, pages 1–5. IEEE, 2026. doi: 10.1109/ICCIDS69108.2026.11407681. URL [References/Semantic_Intent_Weaver_A_Context-Aware_Search_Algorithm_for_Ambiguous_Query_Resolution.pdf](#).
- [37] Prakash M. Nadkarni, Lucila Ohno-Machado, and Wendy W. Chapman. Natural language processing: an introduction. *Journal of the American Medical Informatics Association*, 18(5): 544–551, 2011. doi: 10.1136/amiajnl-2011-000464.
- [38] Benjamin Shickel, Patrick James Tighe, and Azra Bihorac. Deep ehr: A survey of recent advances in deep learning techniques for electronic health record (ehr) analysis. *IEEE Journal of Biomedical and Health Informatics*, 2017.
- [39] Edward Choi, Mohammad Taha Bahadori, Andy Schuetz, Walter F. Stewart, and Jimeng Sun. Doctor ai: Predicting clinical events via recurrent neural networks. In *Machine Learning for Healthcare*, 2016.
- [40] Riccardo Miotto, Li Li, Brian A. Kidd, and Joel T. Dudley. Deep patient: An unsupervised representation to predict the future of patients from the electronic health records. *Scientific Reports*, 6(1):26094, May 2016. ISSN 2045-2322. doi: 10.1038/srep26094. URL <https://doi.org/10.1038/srep26094>.
- [41] Alvin Rajkomar, Eyal Oren, Kai Chen, Andrew M. Dai, Nissan Hajaj, Michaela Hardt, Peter J. Liu, Xiaobing Liu, Jake Marcus, Mimi Sun, Patrik Sundberg, Hector Yee, Kun Zhang, Yi Zhang, Gerardo Flores, Gavin E. Duggan, Jamie Irvine, Quoc Le, Kurt Litsch, Alexander Mossin, Justin Tansuwan, De Wang, James Wexler, Jimbo Wilson, Dana Ludwig, Samuel L. Volchenbourn, Katherine Chou, Michael Pearson, Srinivasan Madabushi, Nigam H. Shah, Atul J. Butte, Michael D. Howell, Claire Cui, Greg S. Corrado, and Jeffrey Dean. Scalable and accurate deep learning with electronic health records. *npj Digital Medicine*, 1(1):18, May 2018. ISSN 2398-6352. doi: 10.1038/s41746-018-0029-1. URL <https://doi.org/10.1038/s41746-018-0029-1>.
- [42] Yikuan Li, Shishir Rao, José Roberto Ayala Solares, Abdelaali Hassaine, Remy Ramakrishnan, Dexter Canoy, Yajie Zhu, Kazem Rahimi, and Gholamreza Salimi-Khorshidi. Behrt: Transformer for electronic health records. *Scientific Reports*, 10:7155, 2020.

- [43] Laila Rasmy, Yang Xiang, Zhipeng Xie, Cui Tao, and Degui Zhi. Med-bert: Pretrained contextualized embeddings on large-scale structured electronic health records for disease prediction. *npj Digital Medicine*, 4:86, 2021.
- [44] Junyuan Shang, Tengfei Ma, Cao Xiao, and Jimeng Sun. Pre-training of graph augmented transformers for medication recommendation. In *International Joint Conference on Artificial Intelligence (IJCAI)*, 2019.
- [45] Kexin Huang, Jaan Altosaar, and Rajesh Ranganath. Clinicalbert: Modeling clinical notes and predicting hospital readmission. arXiv preprint, 2019.
- [46] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. Biobert: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 2020. doi: 10.1093/bioinformatics/btz682.
- [47] Emily Alsentzer, John Murphy, William Boag, Wei-Hung Weng, Di Jindi, Tristan Naumann, and Matthew McDermott. Publicly available clinical bert embeddings. In *Proceedings of Clinical NLP*, 2019.
- [48] Lavender Yao Jiang, Xujin Liu, Nima Tsai, Allison Furth, Akhil Vaid, Anand Bhargava, and et al. Health system-scale language models are all-purpose prediction engines. *Nature*, 619(7969):357–362, 2023.
- [49] Xi Yang, Aokun Chen, Nima PourNejatian, Hoo Chang Shin, Kaleb E. Smith, Christopher Parisien, Colin Compas, Cheryl Martin, Anthony B. Costa, Mona G. Flores, Ying Zhang, Tanja Magoc, Christopher A. Harle, Gloria Lipori, Duane A. Mitchell, William R. Hogan, Elizabeth A. Shenkman, Jiang Bian, and Yonghui Wu. A large language model for electronic health records. *npj Digital Medicine*, 5:194, 2022.
- [50] Karan Singhal, Shekoofeh Azizi, Tao Tu, S. Sara Mahdavi, Jason Wei, Hyung Won Chung, Nathan Scales, and et al. Large language models encode clinical knowledge. *Nature*, 620(7972):172–180, 2023.
- [51] Taeyoon Kwon, Wonjin Lee, and Jaewoo Kang. Large language models are clinical reasoners: Reasoning-aware diagnosis via prompting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pages 13210–13218, 2024.
- [52] Adibvafa Fallahpour, Mahshid Alinoori, Wenqian Ye, Xu Cao, Arash Afkanpour, and Amrit Krishnan. Ehrmamba: Towards generalizable and scalable foundation models for electronic health records. In *Proceedings of the 4th Machine Learning for Health Symposium*, volume 259 of *Proceedings of Machine Learning Research*, pages 291–307. PMLR, 2025.
- [53] Lin Lawrence Guo, Adam Paul Yan, Emily Vettese, and Lillian Sung. Porter: Language-grounded event representations for portable structured ehr foundation models, 2026.
- [54] Saeed Shurrah, Mariam Al-Omari, Dana El Samad, and Farah E. Shamout. Ehr-ragp: Retrieval-augmented prototype-guided foundation model for electronic health records, 2026.

- [55] Siru Liu, Allison B. McCoy, and Adam Wright. Improving large language model applications in biomedicine with retrieval-augmented generation: a systematic review, meta-analysis, and clinical development guidelines. *Journal of the American Medical Informatics Association*, 32(4):605–615, 2025. doi: 10.1093/jamia/ocaf008.
- [56] Michael Wornow, Yizhe Xu, Rahul Thapa, Bhavik Patel, Ethan Steinberg, Scott Fleming, Michael A. Pfeffer, Jason Fries, and Nigam H. Shah. The shaky foundations of large language models and foundation models for electronic health records. *npj Digital Medicine*, 6:135, 2023.
- [57] John Sweller. Cognitive load during problem solving: Effects on learning. *Cognitive Science*, 12(2):257–285, 1988.
- [58] Mark A. Musen, Blackford Middleton, and Robert A. Greenes. Clinical decision-support systems. In *Biomedical Informatics*, pages 643–674. Springer, London, 2014.
- [59] Reed T. Sutton, David Pincock, Daniel C. Baumgart, Daniel C. Sadowski, Richard N. Fedorak, and Karen I. Kroeker. An overview of clinical decision support systems: Benefits, risks, and strategies for success. *NPJ Digital Medicine*, 3:17, 2020.
- [60] Yinghao Zhu, Dehao Sui, Zixiang Wang, Xuning Hu, Lei Gu, Yifan Qi, Tianchen Wu, Ling Wang, Yuan Wei, Wen Tang, Zhihan Cui, Yasha Wang, Lequan Yu, Ewen M. Harrison, Junyi Gao, and Liantao Ma. Augmenting clinical decision-making with an interactive and interpretable ai copilot: A real-world user study with clinicians in nephrology and obstetrics. In *Proceedings of the 2026 CHI Conference on Human Factors in Computing Systems*. ACM, 2026.
- [61] David Moher, Alessandro Liberati, Jennifer Tetzlaff, and Douglas G Altman. Preferred reporting items for systematic reviews and meta-analyses: the prisma statement. *bmj*, 2009.
- [62] Kenneth F. Schulz, Douglas G. Altman, and David Moher. Consort 2010 statement: updated guidelines for reporting parallel group randomised trials. *BMC Medicine*, 8(1):18, 2010. doi: 10.1186/1741-7015-8-18.
- [63] John Brooke. Sus: A ”quick and dirty” usability scale. In *Usability Evaluation in Industry*, pages 189–194. Taylor & Francis, London, 1996.
- [64] Sandra G. Hart and Lowell E. Staveland. Development of nasa-tlx (task load index): Results of empirical and theoretical research. In *Advances in Psychology*, volume 52, pages 139–183. Elsevier, 1988.
- [65] Joshua C. Mandel, David A. Kreda, Kenneth D. Mandl, Isaac S. Kohane, and Rachel B. Ramoni. Smart on fhir: a standards-based, interoperable apps platform for electronic health records. *Journal of the American Medical Informatics Association*, 23(5):899–908, 2016.
- [66] Karim Lekadir, Ben Glocker, Celia Cintas, Curtis P Langlotz, Eva Weicken, Fred Prior, Georgios Kaissis, Gianna Tsakou, Jayashree Kalpathy-Cramer, John Mongan, Kaisar Kushibar,

- Katrine Riklund, Kostas Marias, and AI training. Future-ai: international consensus guideline for trustworthy and deployable artificial intelligence in healthcare. *BMJ*, 2025.
- [67] Feng Chen, Liqin Wang, Julie Hong, Jiaqi Jiang, and Li Zhou. Unmasking bias in artificial intelligence: a systematic review of bias detection and mitigation strategies in electronic health record-based models. *Journal of the American Medical Informatics Association*, 2024.
- [68] Cecilia S. Lee and Aaron Y. Lee. Clinical applications of continual learning machine learning models. *The Lancet Digital Health*, 2(6):e279–e281, 2020. doi: 10.1016/S2589-7500(20)30102-3.
- [69] Julia Amann, Alessandro Blasimme, Effy Vayena, Dietmar Frey, Vince I. Madai, and the Precise4Q consortium. Explainability for artificial intelligence in healthcare: a multidisciplinary perspective. *BMC Medical Informatics and Decision Making*, 20(1):310, November 2020. ISSN 1472-6947. doi: 10.1186/s12911-020-01332-6. URL <https://doi.org/10.1186/s12911-020-01332-6>.
- [70] Stephen E. Robertson, Steve Walker, Susan Jones, Micheline M. Hancock-Beaulieu, and Mike Gatford. Okapi at trec-3. In *Proceedings of the Third Text REtrieval Conference (TREC-3)*, 1995.
- [71] Gerard Salton and Michael J. McGill. *Introduction to Modern Information Retrieval*. McGraw-Hill, 1983.
- [72] Alistair E. W. Johnson, Tom J. Pollard, Lu Shen, Li-wei H. Lehman, Mengling Feng, Mohammad Ghassemi, Benjamin Moody, Peter Szolovits, Leo Anthony Celi, and Roger G. Mark. Mimic-iii, a freely accessible critical care database. *Scientific Data*, 3:160035, 2016. doi: 10.1038/sdata.2016.35.
- [73] Alistair E. W. Johnson, Lucas Bulgarelli, Lu Shen, Alvin Gayles, Ayad Shammout, Steven Horng, Tom J. Pollard, Sicheng Hao, Benjamin Moody, Brian Gow, Li-wei H. Lehman, Leo Anthony Celi, and Roger G. Mark. Mimic-iv, a freely accessible electronic health record dataset. *Scientific Data*, 10(1):1, 2023. doi: 10.1038/s41597-022-01899-x.
- [74] Alistair Johnson, Tom Pollard, Steven Horng, Leo Anthony Celi, and Roger Mark. Mimic-iv-note: Deidentified free-text clinical notes (version 2.2), 2023. URL <https://doi.org/10.13026/1n74-ne17>.
- [75] Alistair E. W. Johnson, Tom J. Pollard, Seth J. Berkowitz, Nathaniel R. Greenbaum, Matthew P. Lungren, Chih-ying Deng, Roger G. Mark, and Steven Horng. Mimic-cxr, a de-identified publicly available database of chest radiographs with free-text reports. *Scientific Data*, 6:317, 2019. doi: 10.1038/s41597-019-0322-0.
- [76] Benjamin Moody, Sicheng Hao, Brian Gow, Tom Pollard, Wei Zong, and Roger Mark. Mimic-iv waveform database (version 0.1.0), 2022. URL <https://doi.org/10.13026/a2mw-f949>.

- [77] Tom J. Pollard, Alistair E. W. Johnson, Jesse D. Raffa, Leo A. Celi, Roger G. Mark, and Omar Badawi. The eicu collaborative research database, a freely available multi-center database for critical care research. *Scientific Data*, 5:180178, 2018. doi: 10.1038/sdata.2018.178.
- [78] Martin Faltys, Marc Zimmermann, Xinrui Lyu, Matthias Hüser, Stephanie Hyland, Gunnar Rätsch, and Tobias Merz. Hirid, a high time-resolution icu dataset (version 1.1.1), 2021.
- [79] Donald Gross, John F. Shortle, James M. Thompson, and Carl M. Harris. *Fundamentals of Queueing Theory*. Wiley, 4th edition, 2008.
- [80] Stephen E. Robertson and Karen Sparck Jones. Relevance weighting of search terms. *Journal of the American Society for Information Science*, 27(3):129–146, 1976.